




**Prof. Dr. Florian Künzner**

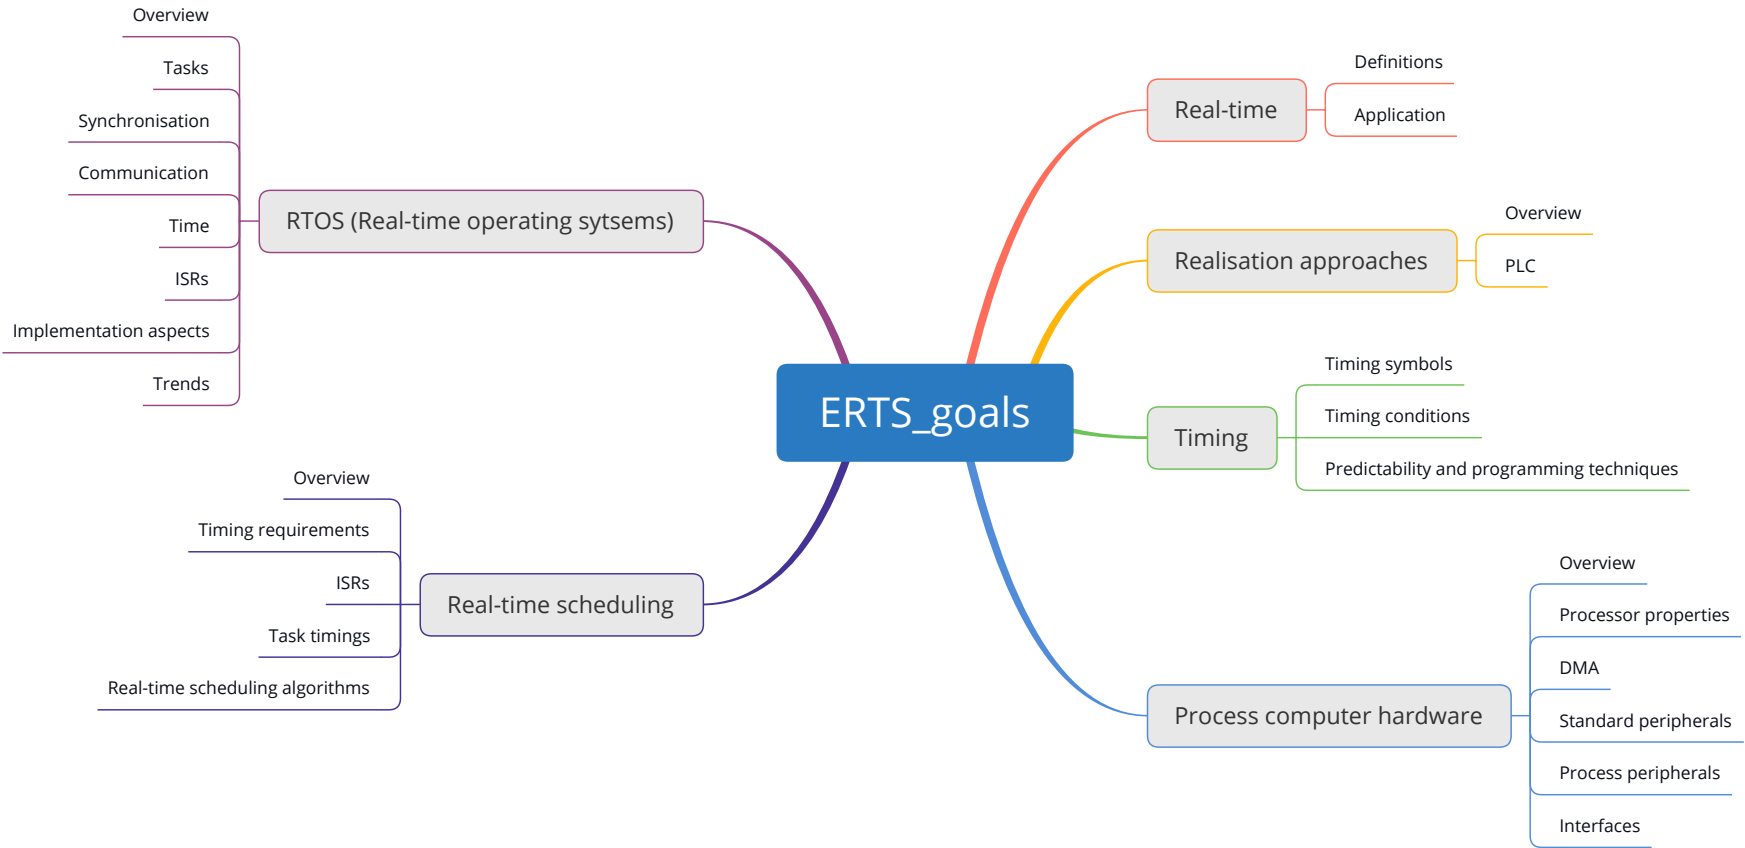
Technical University of Applied Sciences Rosenheim, Computer Science

# ERTS - Embedded real-time systems

**ERTS 4 – Process computer HW 1**



# Goal



# Goal

## ERTS::Process computer hardware 1

- Processor specifics
- Interrupts
- Interrupt interference

# Overview

## Process computer hardware topics:

- Processor specifics
- Interrupts
- Standard peripherals: I/O
- Clocks
- DMA
- Process periphery

## Important but not yet covered in this lecture:

- Field buses: special bus systems, TSN (time sensitive networking)

# Processor specifics

## Architectural features of the hardware (processor + peripherals)

- Pipelining
- Caching
- MMU
- Interrupts
- Direct memory access (DMA)

# Pipelining

Multiple instructions are in the pipeline

1. Fetch	2. Decode	3. Execute	
CMD1			$t_1$
CMD2	CMD1		$t_2$
CMD3	CMD2	CMD1	$t_3$
CMD4	CMD3	CMD2	$t_4$

time  
↓

⇒ speeds-up execution

But: Pipeline flushes and refills for

- jumps/branches
- task changes
- interrupts

⇒ causes **unpredictable runtimes** (variance)

# Caching


## Processor caches (L1, L2, ...)

Key (real adr.)	Value (data: byte 0 to 3)			
	#0	#1	#2	#3
0x0...0100	0x12	0x34	?	?

- ⇒ improves the overall performance
- But: Cache misses occurs (flushes and refills) for
  - jumps/branches
  - task changes
  - interrupts
- ⇒ causes **unpredictable runtimes** (variance)
  - This can increase the WCET by a factor of 33% [2, Buttazzo, p. 15]
  - In a thesis: 10  $\mu$ s measurement accuracy required, but cache introduced 20  $\mu$ s uncertainty

# MMU

The MMU allows paging with all the known security features.



But:

- MMU has to be reprogrammed on process (task) changes
- This introduces some additional processor cycles



# First summary

**Modern processors optimise the overall runtime, but not the WCET.**

## Conclusion:

- Determining the **runtime** very accurate is **extremely hard**
- Adding a reasonable **buffer** to the **WCET**, according to the processor mechanism
- In practice,  $U_{\max} \leq 0.3, \dots, 0.5$  to have sufficient safety reserves.

# Processor

But which processor should we then  
use?

e.g. ARM Cortex-R

# Interrupts

**Hardware interrupts occur asynchronously and are therefore unpredictable.**

**Used for:**

- HW communication
- Controlling the standard peripherals
- Process I/O and alarms

**Details:**

- Different interrupt priorities
- Interrupts preempt application and kernel tasks

# Interrupt handling sequence

## Typical interrupt handling sequence for a CPU:

- 1 CPU accepts interrupt
- 2 Save PC (program counter) and SR (status register); disable interrupts
- 3 Switch to supervisor mode, save registers, load new registers (typically  $< 1\mu s$ )
- 4 Load PC from interrupt vector table
- 5 If there is no HW for 3, then save registers per software (typically  $< 2\mu s$ )
- 6 Execute ISR (interrupt service routine)
- 7 If 5, then restore all register in software
- 8 IRET, restore registers, PC, and SR (return to previous task)

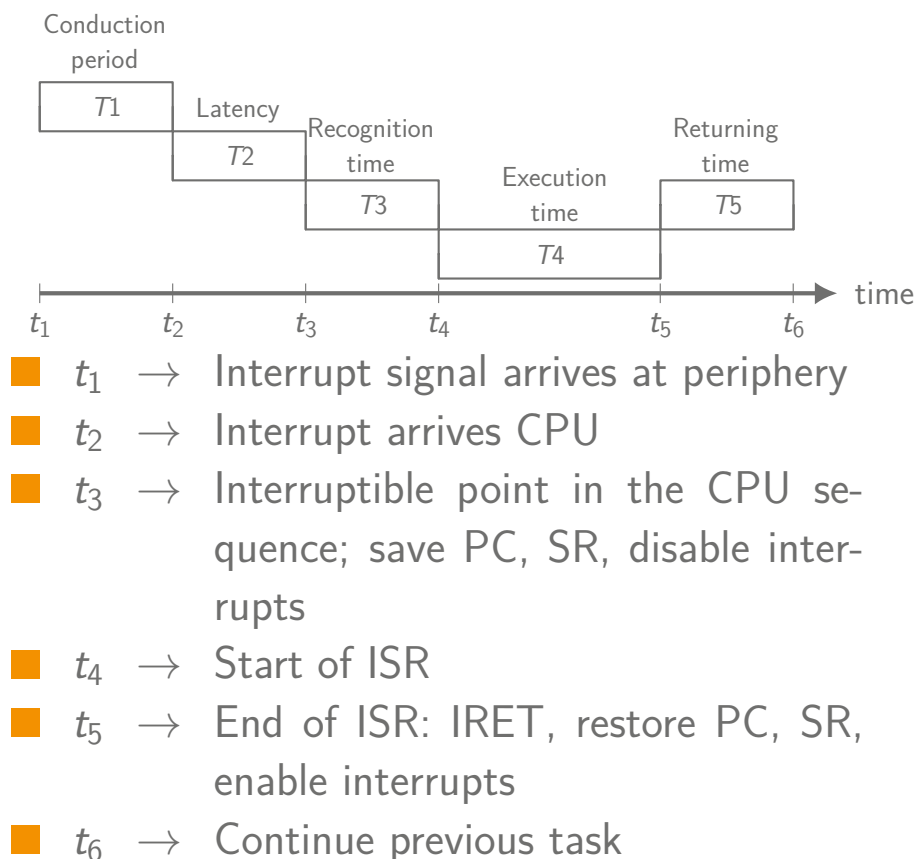
## Discussion:

- Depending on the system, interrupts with a higher priority can interrupt ISRs with a lower priority.
- In all cases, ISRs should be short.



# Interrupt handling timing symbols

## Time characteristics according to DIN 66216\* (page 2)



■  $T_1 \rightarrow$  Conduction period

■  $T_2 \rightarrow$  Latency

■  $T_3 \rightarrow$  Recognition time

■  $T_4 \rightarrow$  Execution time

■  $T_5 \rightarrow$  Returning time

■  $T_{IRE} \rightarrow$  Reaction time:

$$T_{IRE} = T_1 + T_2 + T_3$$

■  $T_{IR} \rightarrow$  Response time:

$$T_{IR} = T_1 + T_2 + T_3 + T_4$$

■  $T_I \rightarrow$  Interrupt time:  $T_I = T_3 + T_4 + T_5$

■  $T_{IO} \rightarrow$  Organisation time (overhead):

$$T_{IO} = T_3 + T_5$$

\*The DIN 66216 is outdated and withdrawn.

# Examples

The interrupt timing information are very **hardware** and **process specific**.

To get a first impression, we look at some examples\*.

---

\*The examples may be outdated, and may need some references.



# Example 1: Typical interrupt rates

- $T_4$  → Execution time, depending on HW (clock rate):  
typically 10...20 instructions:  $T_4 = 0.5 \dots 20 \mu s$
- $T_I$  → Interrupt time:  $T_I = T_3 + T_4 + T_5$   
typically:  $1 \mu s < T_I < 220 \mu s$
- $T_{IO}$  → Organisation time (overhead):  
typically:  $T_{IO} = T_3 + T_5 \Rightarrow T_{IO} = 0.5 \dots 200 \mu s$   
MnP: NXP Coldfire M52259:  $T_{IO} = 2400 ns$
- $R_{I_{max}}$  → Maximum interrupt rate:  $R_{I_{max}}$   
typically:  $4.5 kHz < R_{I_{max}} < 1 MHz$

## Example 2: Reaction time

$T_{IRE}$  depends essentially on the longest non-interruptible phase:

- Longest hardware instruction
- Longest bus usage
- Longest software function with disabled interrupts

Through an OS kernel,  $T_2 + T_3$  usually increases, because OS kernel routines are often (partially) non-interruptible.

**Typical values for  $T_{IRE}$ :**

- $T_{IRE} = 1\text{ms}$  → For bad HW + operating systems combinations
- $T_{IRE} = 0.1, \dots, 20\mu\text{s}$  → For good HW + operating systems combinations
- $T_{IRE} = 1\text{s}$  → For non real-time operating systems (e.g. old standard unix system with exotic services)



# Example 3: Recognition and returning time

**$T_3$  and  $T_5$  depends strongly on the hardware:**

- Which and how many registers are automatically saved/restored
- FPU enabled  $\Rightarrow$  more registers to save/restore?

**$T_3$  example for: MnP: Motorola 68000, 16 bit, 16 MHz:**

- 44 cycles  $\rightarrow 2.75 \mu s$  (without save registers)

$$T_3 = \frac{44}{16 \text{ MHz}} = \frac{44}{16 \times 10^6} = 2.75 \mu s$$

- 84 cycles  $\rightarrow 4.25 \mu s$  (save 7 registers to stack)

84 cycles = 44 cycles + (12 + 4\*7 cycles MOVEM)

$$T_3 = \frac{84}{16 \text{ MHz}} = \frac{84}{16 \times 10^6} = 4.25 \mu s$$

## Example 4: $T1 + T2$

Measured value for  $T1 + T2$  on a ColdFire with 80 MHz:

- $T1 + T2 \rightarrow T1 + T2 = 400, \dots 800\text{ns}$

Data sheet value for  $T1 + T2$  on a NXP: QorIQ P2020, 1200 MHz,  
RTOS (Integrity, Green Hills)

- $T1 + T2 \rightarrow T1 + T2 = 0.144\mu\text{s}$

# Interrupt interference

## Techniques to reduce interrupt interference

# Interrupt interference

## Timing conditions with interrupts

If time conditions can't be met with interrupts (ISRs), the following methods are possible:

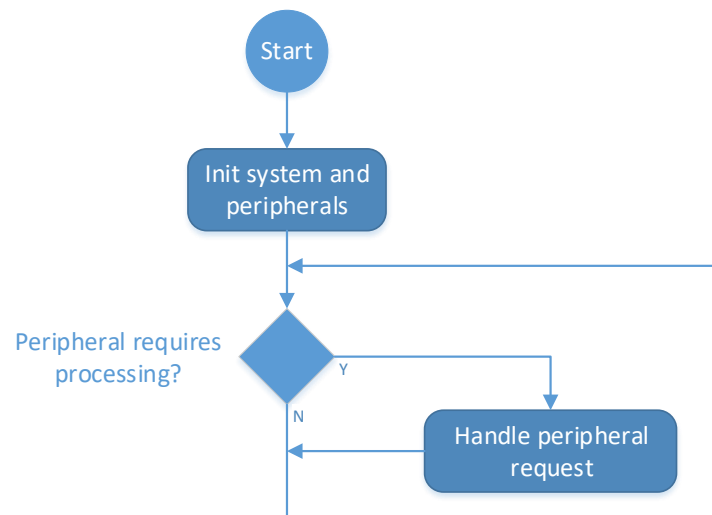
- Reduce interrupt interference:
  - Approach A: Programmed I/O: busy wait or polling
  - Approach B: Periodic hardware tasks
  - Approach C: Delayed interrupt handling (e.g. with RTOS tasks)
- DMA
- Additional or other hardware



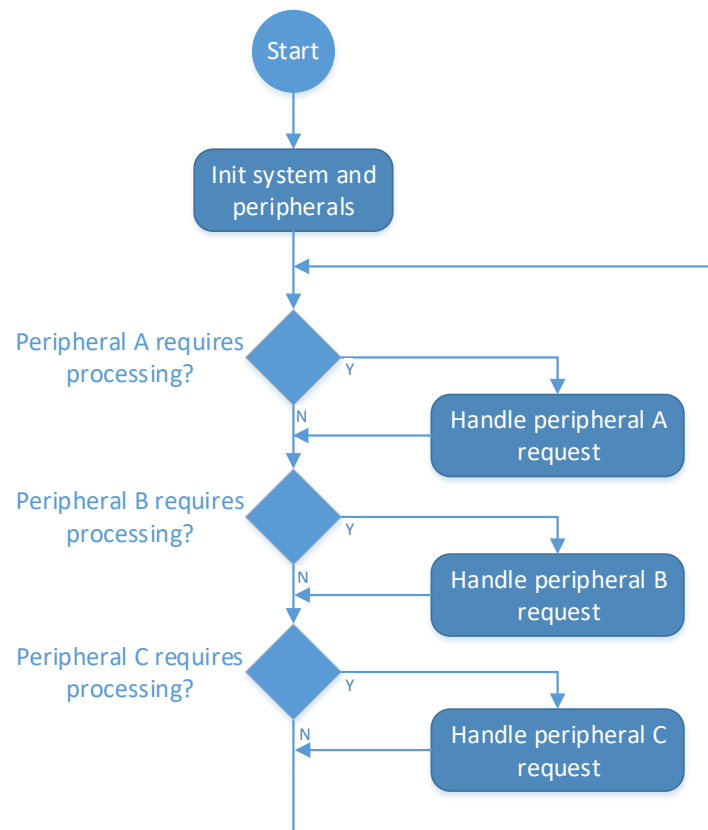
# Interrupt interference: Approach A

## Programmed I/O: busy wait or polling

One peripheral:



Multiple peripherals:





# Interrupt interference: Approach A

## Programmed I/O: busy wait or polling

### Disable all external interrupts

#### Details:

- Disable all external interrupts (except the timer or systick interrupt)
- Access the hardware directly and perform programmed I/O with busy wait or polling

#### Pro:

- Introduces a high predictability

#### Con:

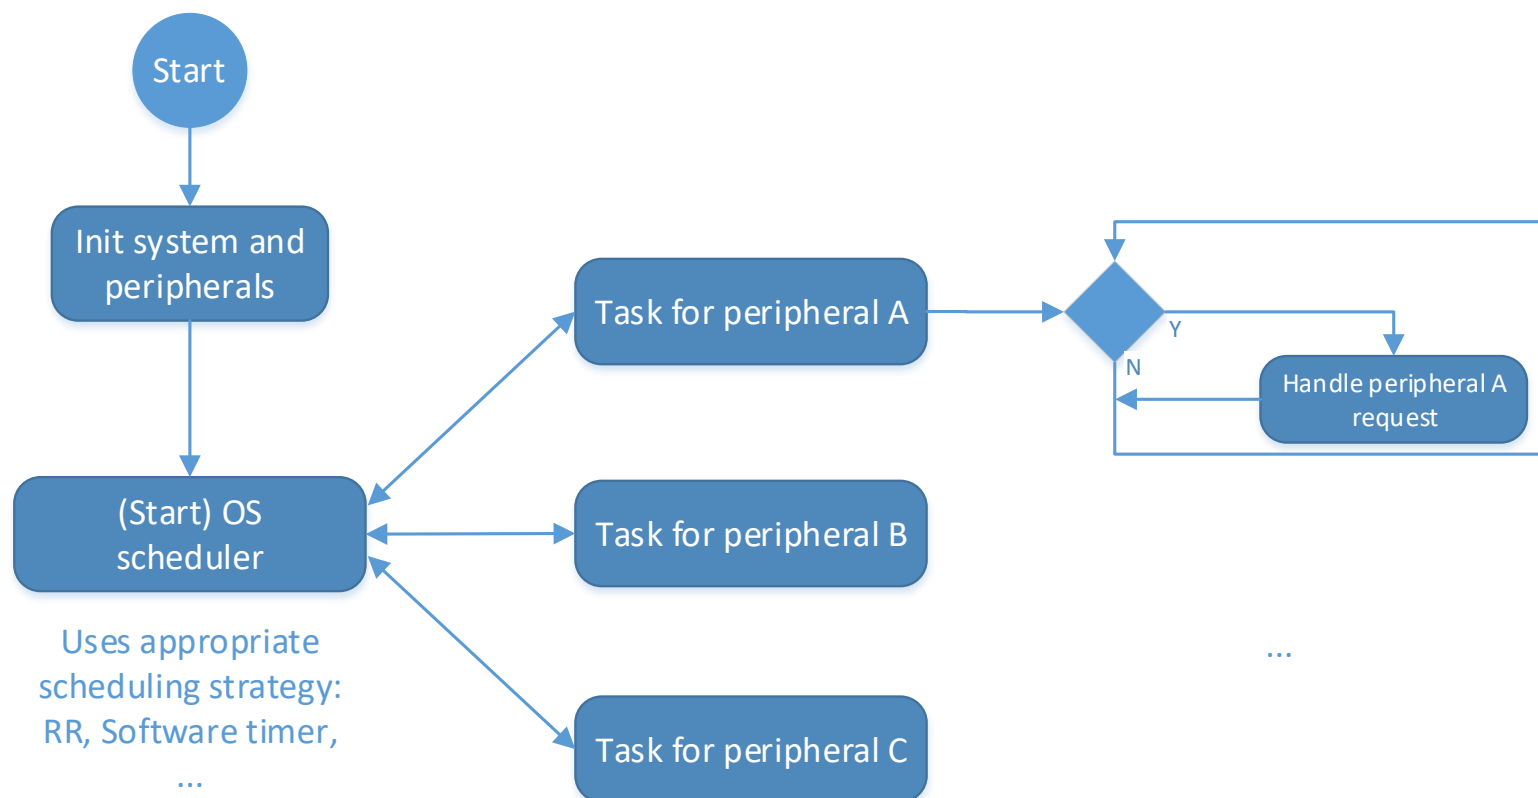
- Low processor efficiency (time spent for busy wait or polling)
- Only possible to poll a few (or one) devices

source: [2, Buttazzo, p. 16–17]



# Interrupt interference: Approach B

## Periodic hardware tasks



# Interrupt interference: Approach B

## Periodic hardware tasks

Use periodic hardware tasks, while external interrupts are disabled

### Details:

- Disable all external interrupts (except the timer or systick interrupt)
- Access the hardware with periodic tasks
  - Tasks can be part of the kernel
  - Different tasks for slow/fast devices

### Pro:

- Hardware details are encapsulated within kernel tasks

### Con:

- Low processor efficiency (time spent for busy wait or polling)

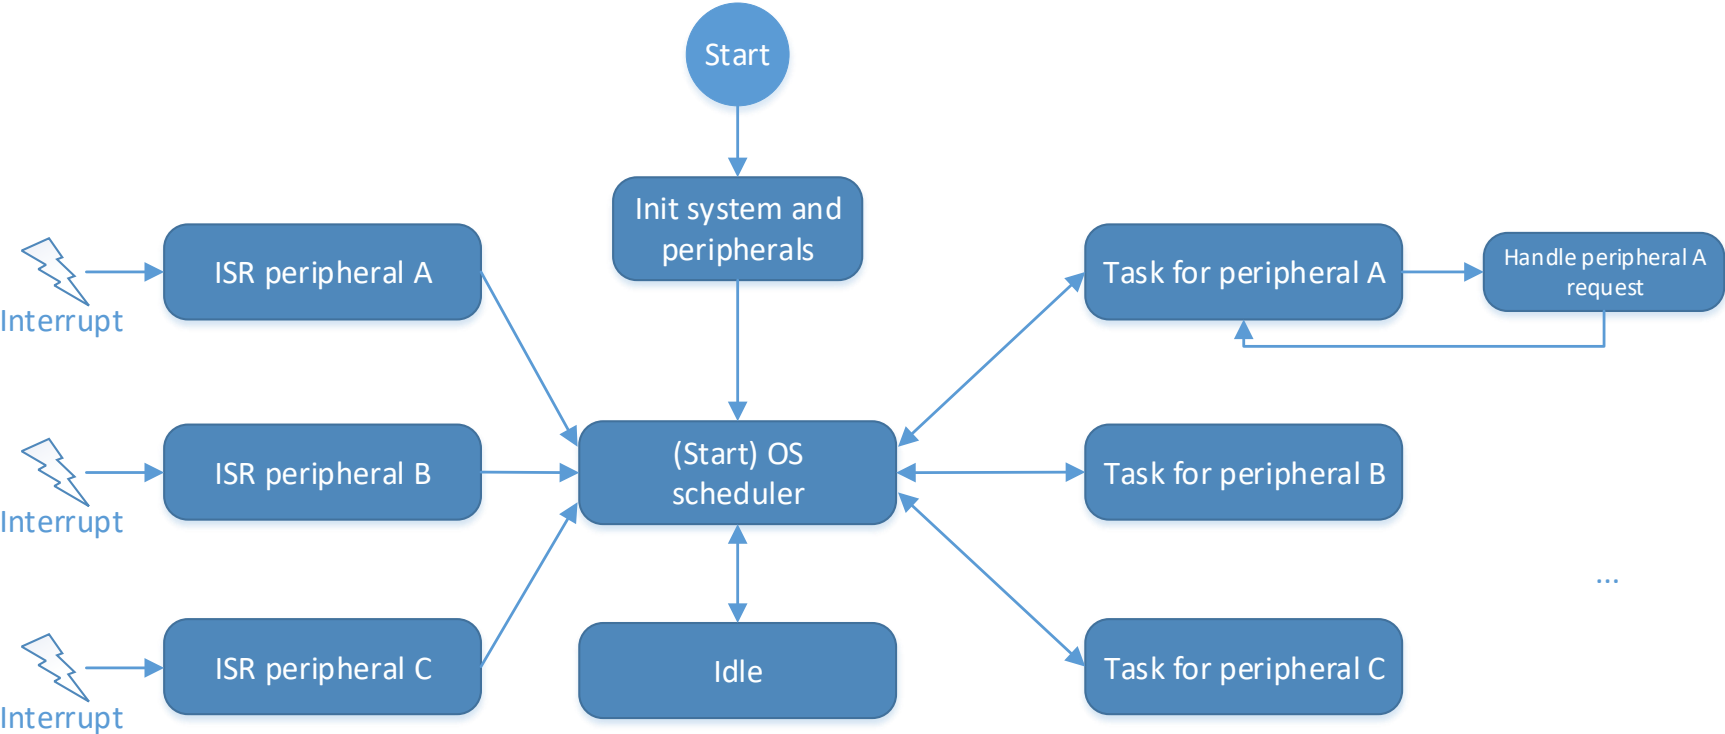
source: [2, Buttazzo, p. 17]

used in: [DRSK89] A. Damm, J. Reisinger, W. Schwabl, and H. Kopetz. The real-time operating system of MARS. Operating System Review, 23(3):141–157, July 1989.



# Interrupt interference: Approach C

## Delayed interrupt handling (e.g. with RTOS tasks)





# Interrupt interference: Approach C

## Delayed interrupt handling (e.g. with RTOS tasks)

### Use ISR to handling task activation

#### Details:

- All external interrupts kept enabled
- The ISR is very short and only activates a handling task
- The handling task manages the device

#### Pro:

- Eliminates busy wait and polling overhead
- Unbounded delays introduced by driver tasks are drastically eliminated
- A control task can have a higher priority than a device handling task.

#### Con:

- Still contains some unbounded delays for the ISRs

source: [2, Buttazzo, p. 17–18]

used in: [TM89] H. Tokuda and C. W. Mercer. ARTS: A distributed real-time kernel. Operating System Review, 23(3), July 1989.

# Exercise

Try to find as much timing ( $T1 \dots T5$ ) information as possible for the Cortex M4 architecture!

## References:

- Cortex-M4 Technical Reference Manual:  
<https://developer.arm.com/documentation/ddi0439/b/>

# Exercise

## Cortex M4 architecture: Interrupt handling timing

- $T_1$  → Conduction period:  $T_1 =$
- $T_2$  → Latency:  $T_2 =$
- $T_3$  → Recognition time:  $T_3 =$
- $T_4$  → Execution time:  $T_4 =$
- $T_5$  → Returning time:  $T_5 =$
- $T_{IRE}$  → Reaction time:  
$$T_{IRE} = T_1 + T_2 + T_3 =$$
- $T_{IR}$  → Response time:  
$$T_{IR} = T_1 + T_2 + T_3 + T_4 =$$
- $T_I$  → Interrupt time:  $T_I = T_3 + T_4 + T_5 =$
- $T_{IO}$  → Organisation time (overhead):  $T_{IO} = T_3 + T_5 =$

# Summary and outlook

## Summary

- Processor specifics
- Interrupts
- Interrupt interference

## Outlook

- Digital and analog I/O
- Real-time clocks