



**Prof. Dr. Florian Künzner**

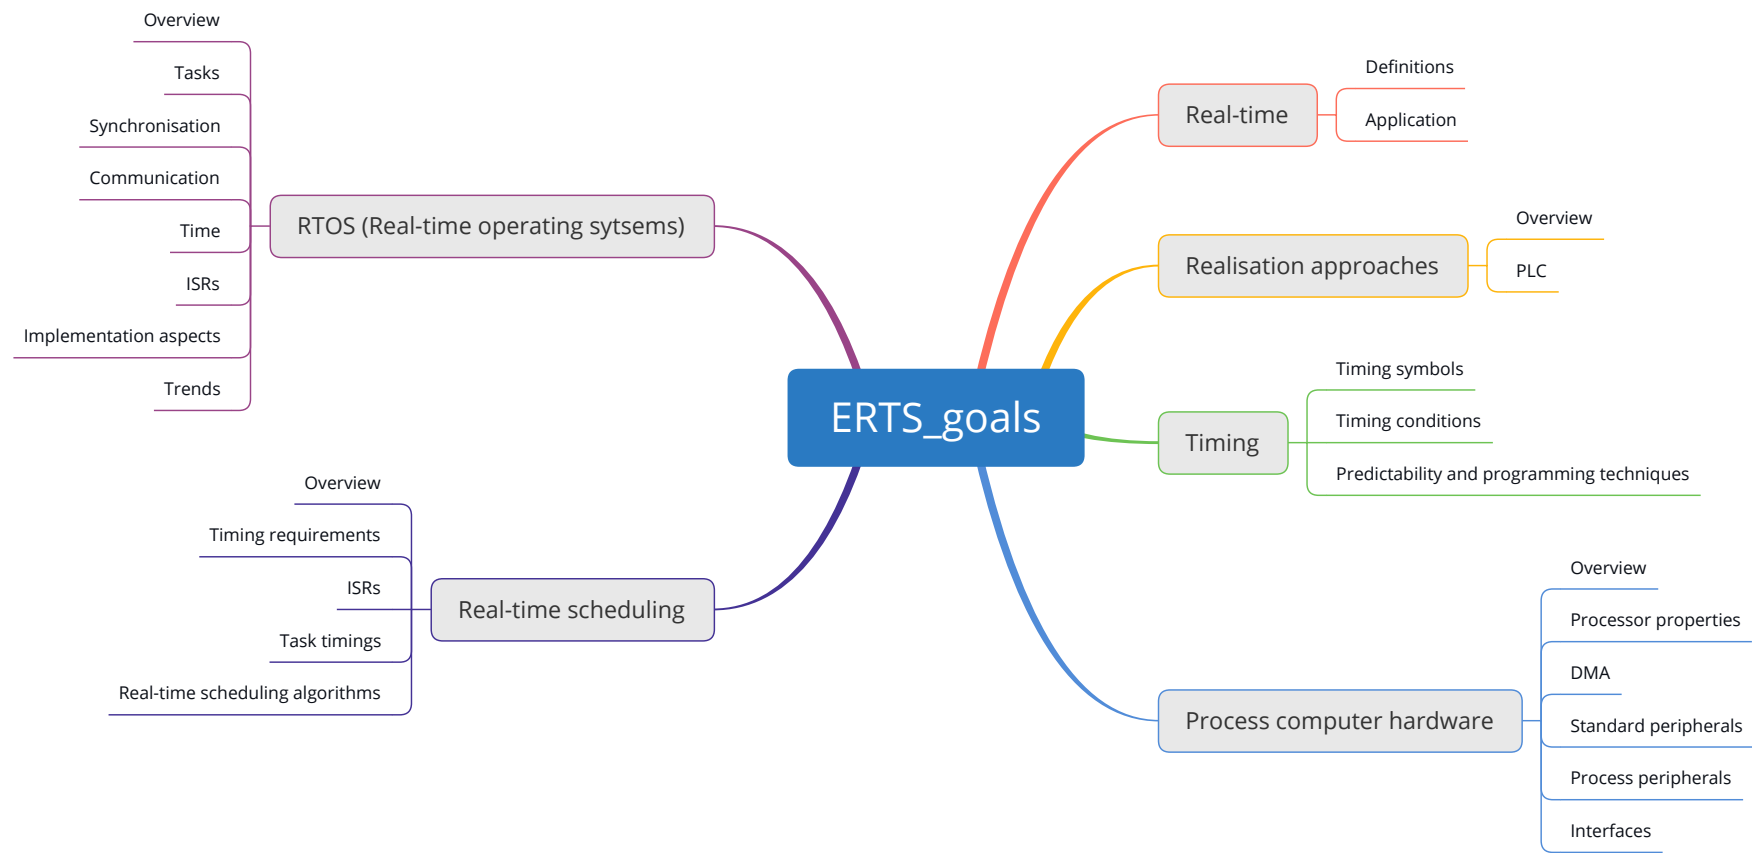
Technical University of Applied Sciences Rosenheim, Computer Science

# ERTS - Embedded real-time systems

**ERTS 6 – Process computer HW 3**



# Goal





# Goal

## ERTS::Process computer HW 3

- Overview
- DMA interference
- DMA system load examples
- DMA system load definitions
- DMA Exercises

# Overview

## How to transfer data between devices and the RAM?



# Overview

## I/O programming methods\*:

- **Programmed I/O** (busy wait or polling)
  - CON: CPU has to check if device is ready (for every word transfer)
  - CON: CPU has to perform the data transfer
- **Interrupt driven I/O**
  - + PRO: CPU is informed when device is ready (for every word transfer)
  - CON: CPU has to perform the data transfer
- **Direct memory access (DMA)**
  - + PRO: BUS devices directly transfer data
  - + PRO: CPU is informed after everything has been finished

---

\*recap from RA lecture

# DMA recap

Recap from **CA** lecture.

# Interrupt driven I/O

System load example for  
**interrupt driven I/O**

# (a) Interrupt driven I/O (single word transfer)

## Given:

- I/O transfer rate:  $R_{TIO} = 64000$  words/s
- Bus width: 1 word = 16 bit
- To transfer: 1 word, 10 instructions (1 inst.  $\approx$  1 cycle  $\rightarrow$  10 cycles) are required for the ISR
- The processor runs with 1 MHz

## Question:

- What is the resulting processor utilisation  $U_P$  (also known as system load)?
- Transfer rate  $R_T$  in kB/s?

## Proposed solution:

- $T_E = 10 \times \frac{1}{1 \text{ MHz}} = 10 \times \frac{1}{1 \times 10^6} = 0.00001 \text{ s} \Rightarrow 10 \text{ } \mu\text{s}$
- $T_P = \frac{1}{64000} = 0,000015625 \text{ s} \Rightarrow 15.625 \text{ } \mu\text{s}$
- $U_P = \frac{T_E}{T_P} = \frac{10 \text{ } \mu\text{s}}{15.625 \text{ } \mu\text{s}} = 0.64 \Rightarrow 64 \%$
- Transfer rate:  $R_T = 64000 \times \frac{16}{8} = 128000 \text{ B} \Rightarrow 128 \text{ kB/s}$



# Interrupt driven I/O

Can **DMA** improve the situation?



# DMA interference

## *Cycle stealing method (while the DMA device is active):*

- Both, the I/O transfer and the CPU runs in parallel
- The CPU can't use the memory
- If the CPU needs the memory (read/write), the CPU has to wait until the DMA cycle is completed

### **Problem:**

- No way of predicting how many times the CPU will have to wait for the DMA cycle completion

⇒ Response time of running task cannot be precisely determined

## **Another method: *time-slice method* [SR88]**

- Each memory cycle is split into two adjacent time slots
- One for the CPU, the other for the DMA device
- + Pro: More predictable, due to no or less CPU waiting conditions
- Con: More expensive than the cycle stealing

source: [2, Buttazzo, p. 14]

[SR88] J. Stankovic and K. Ramamritham, editors. Tutorial on Hard Real-Time Systems. IEEE Computer Society Press, 1988.



# DMA: single word transfer vs burst mode

For the **DMA**, a distinction is made between **single word transfers** and a **burst mode** with  $b_{max} = 1, \dots, \infty$ .

*Hint:* **Large**  $b_{max}$  are **problematic** for real-time systems, because it can reduce the reactivity and introduce some **non-determinism** in **predicting** the **WCET**.

# DMA system load definitions

What is the influence of the **DMA**: Load or bus utilisation and the increased, resulting runtime?

# (a) DMA (single word transfer)

## Resulting system load or bus utilisation through DMA:

- $\beta$  or  $U_{\text{BUS}}$  → Load or bus utilisation  $\beta = n \times T_{\text{DMA}_c}$
- $n$  → Number of words per second
- $T_{\text{DMA}_c}$  → Time for one DMA cycle

## Details:

- Typical values for  $T_{\text{DMA}_c} = 500, \dots, 100$  ns for slower systems (VMEBus)
- For normal PCs,  $T_{\text{DMA}_c}$  is faster

## Resulting runtime:

- There is a maximum approximated increase of the runtime by the factor:  $(1 + \beta)$
- Resulting runtime  $\approx$  runtime for a calculation or a task  $\times (1 + \beta)$
- Attention: This is only valid for small  $n$ ; the exact factor is  $\frac{1}{1-\beta}$

## (b) DMA (burst mode, fixed block size)

### Resulting system load or bus utilisation through DMA:

- $\beta$  or  $U_{\text{BUS}}$   $\rightarrow$  Load or bus utilisation  $\beta = n_b \times T_{\text{DMA}_b}$
- $n_b$   $\rightarrow$  Number of burst blocks per second
- $T_{\text{DMA}_b}$   $\rightarrow$  Time to transfer one burst block:  $T_{\text{DMA}_b} = T_{\text{latency}} + b_{\text{max}} \times T_{\text{DMA}_c}$
- $T_{\text{latency}}$   $\rightarrow$  Waiting time after a burst block

### Resulting runtime:

- There is a maximum approximated increase of the runtime by the factor:  $\frac{1}{1-\beta}$
- Resulting runtime  $\approx$  runtime for a calculation or a task  $\times \frac{1}{1-\beta}$

# DMA system load examples

## Detailed analysis:

- DMA (single word transfer, without cache, for theoretical purposes only)
- DMA (burst mode, normal load)
- DMA (extreme/high load)

## (b) DMA (single word transfer)

Details: without cache, for theoretical purposes only

### Given:

- To transfer:  $R_{TIO} = 64000$  words/s
- Bus width: 1 word = 16 bit
- A/D converter (ADC) generates 1 word every  $15.625 \mu s$
- ⇒ Rate:  $R_{ADC} = \frac{1}{15.625 \mu s} = 64000$  samples/s
- Bus cycle time: 150 ns

### Question:

- What is the resulting load or bus utilisation  $\beta$  through the cycle stealing?
- What is the maximum transfer rate  $R_{T_{BUS}}$  of the bus?

### Proposed solution:

- Load:  $\beta = 64000 \times 150 \text{ ns} = 0.0096 \approx 0.01 \Rightarrow 1 \%$
- Maximum transfer rate of bus:  

$$R_{T_{BUS}} = \frac{1}{150 \text{ ns}} \times \frac{16}{8} = \frac{1}{150 \times 10^{-9}} \times 2 = 13333333.33 \text{ B} \Rightarrow \approx 13.33 \text{ MB/s}$$

**But:** Modern systems are more complex!



## (c) DMA (burst mode, normal load)

### Given:

- To transfer: 64000 words/s, each word has again 16 bit.
- A/D converter (ADC) generates 1 word every  $15.625 \mu\text{s}$ ;  $R_{ADC} = \frac{1}{15.625 \mu\text{s}} = 64000$  samples
- Bus width: 1 word = 64 bit
- Transfer time: 2.5 ns for one word;  $T_{\text{latency}} = 50$  ns (after a block transfer)
- $b_{\text{max}} = 8$ ; all memory accesses and transfers are 64 byte blocks

### Question:

- What is the resulting load or bus utilisation  $\beta$  through the cycle stealing?

### Proposed solution:

- Number of burst blocks per second:  $n_b = \frac{64000}{8} = 8000$
- Time to transfer one burst block:  $T_{\text{DMA}_b} = 50 \text{ ns} + 8 \times 2.5 \text{ ns} = 70 \text{ ns}$
- Load:  $\beta = 8000 \times 70 \text{ ns} = 0.00056 \Rightarrow 0.056 \%$

## (d) DMA (burst mode, high load)

### Given:

- To transfer:  $1000 \times 64000$  words/s = 64,000,000 words/s, each word has again 16 bit.
- $R_{ADC} = 1000 \times 64000$  samples
- Bus width: 1 word = 64 bit
- Transfer time: 2.5 ns for one word;  $T_{\text{latency}} = 50$  ns (after a block transfer)
- $b_{\text{max}} = 8$ ; all memory accesses and transfers are 64 byte blocks

### Question:

- What is the resulting load or bus utilisation  $\beta$  through the cycle stealing?

### Proposed solution:

- Number of burst blocks per second:  $n_b = \frac{1000 \times 64000}{8} = 80,000,000$
- Load:  $\beta = 80,000,000 \times 70 \text{ ns} = 0.56 \Rightarrow 56 \%$

# Exercise 1: Processor utilisation

## Evaluation of an interrupt driven I/O system

### Given:

- I/O transfer rate:  $R_{TIO} = 40000$  words/s
- Bus width: 1 word = 32 bit
- ISR takes  $T_E = 20 \mu s$

### Question:

- Is it possible to safely handle the load of the I/O? *Hint: You may calculate  $T_P$  and  $U_P$  to answer that.*
- Will interrupts be lost?
- Also consider the influence of the operating system!




# Exercise 1: Processor utilisation

## Evaluation of an interrupt driven I/O system

### Proposed solution:

- $T_P = \frac{1}{40000} = 0.000025 \text{ s} \Rightarrow 25 \mu\text{s}$
- $U_P = \frac{T_E}{T_P} = \frac{20 \mu\text{s}}{25 \mu\text{s}} = 0.8 \Rightarrow 80 \%$

### Interrupt period analysis:



- (1) If the interrupts are equally distributed (small variance/jitter), then this might be possible
- (2) Here, it's very close with no safety buffer
- (3) Here, interrupt signals are lost

### Evaluation:

- On a hard real-time system (and safety), usually, (1), (2), and (3) are not acceptable
- For a soft or firm real-time system, (1) may be acceptable (but usually not)
- If an OS is used and if it disables the interrupt for some time, this might further complicate the situation

# Exercise 2: DMA system load

## Evaluation of an I/O system using DMA

### Given:

- I/O transfer: 10 MiB within 1 s
- Bus width: 1 word = 32 bit
- Block length = 64 bytes

### Question:

- Is it possible to safely handle the load of the I/O? *Hint: You may calculate the load  $\beta$  to answer that.*
- Will interrupts be lost?
- Also consider the influence of the operating system!

# Exercise 2: DMA system load

## Evaluation of an I/O system using DMA

### Proposed solution:

- 1 word:  $\frac{32}{8} = 4$  byte;  $b_{max} = \frac{64}{4} = 16$
- $n_b = \frac{10 \text{ MiB}}{64 \text{ bytes}} = \frac{10485760}{64} = 163840$  burst blocks per second
- $T_{DMA_b} = T_{latency} + b_{max} \times T_{DMA_c} = 50 \text{ ns} + 16 \times 2.5 \text{ ns} = 90 \text{ ns}$
- $\beta = n_b \times T_{DMA_b} = 163840 \times 90 \text{ ns} = 0.0147456 \approx 0.015 \Rightarrow 1.5 \%$
- Maximum approximated increase of the runtime by the factor:  $\frac{1}{1-0.015} \approx 1.015$

### Evaluation:

- The increase of the runtime by the factor of  $\approx 1.015$  is low and can usually be accepted for soft, firm, and hard real-time systems
- Also, an RTOS with a limited time where the interrupts are disabled should be acceptable

# Exercise 1+2: System load

## Evaluation of an combined interrupt driven I/O and DMA system

### Question:

- Can the load of example 1 and 2 be combined?

### Proposed solution:

- Estimation:  $U \approx U_P + \beta = 0.8 + 0.015 = 0.815$
- $T_E = 20 \mu s * 1.015 \approx 20.3 \mu s$
- $U = \frac{20.3 \mu s}{25 \mu s} = 0.812 \Rightarrow 81.2 \%$

### Evaluation:

- The additional load introduced by DMA plays practically no role
- But, as stated in example 1, the load of  $\gtrsim 80\%$  is usually not acceptable

# Summary and outlook

## Summary

- Overview
- DMA interference
- DMA system load examples
- DMA system load definitions
- DMA Exercises

## Outlook

- Scheduling theory