



Prof. Dr. Florian Künzner

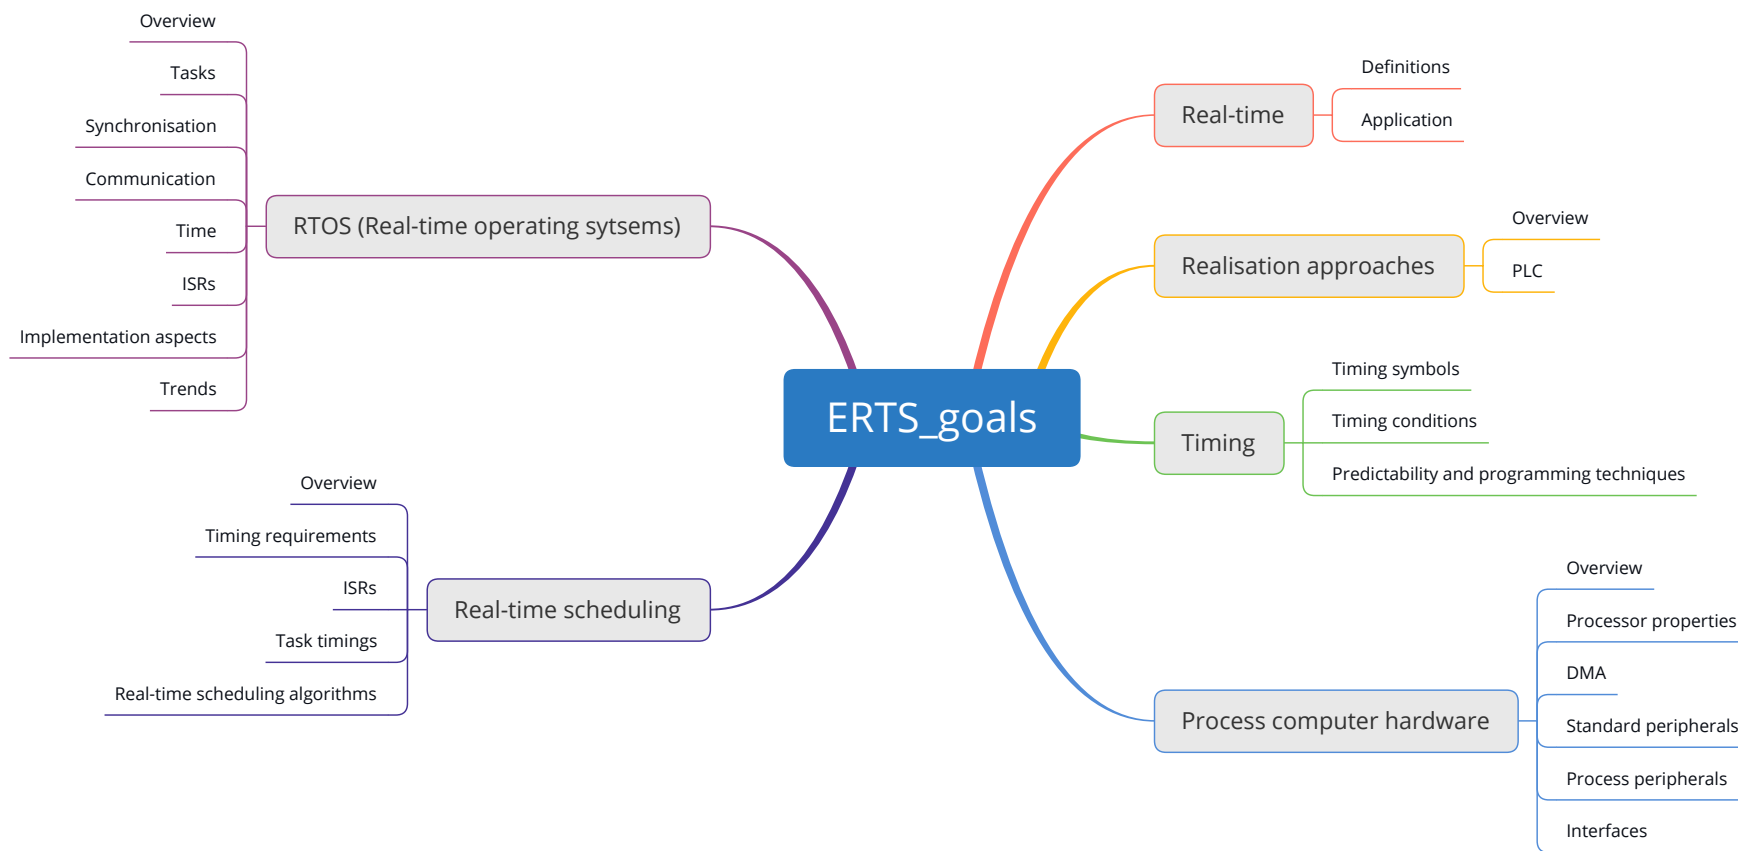
Technical University of Applied Sciences Rosenheim, Computer Science

ERTS - Embedded real-time systems

ERTS 8 – RTOS



Goal



Goal

ERTS::Real-time operating systems

- Intro to RTOS
- Tasks
- Synchronisation
- Communication
- Interrupts
- Timers
- Trends



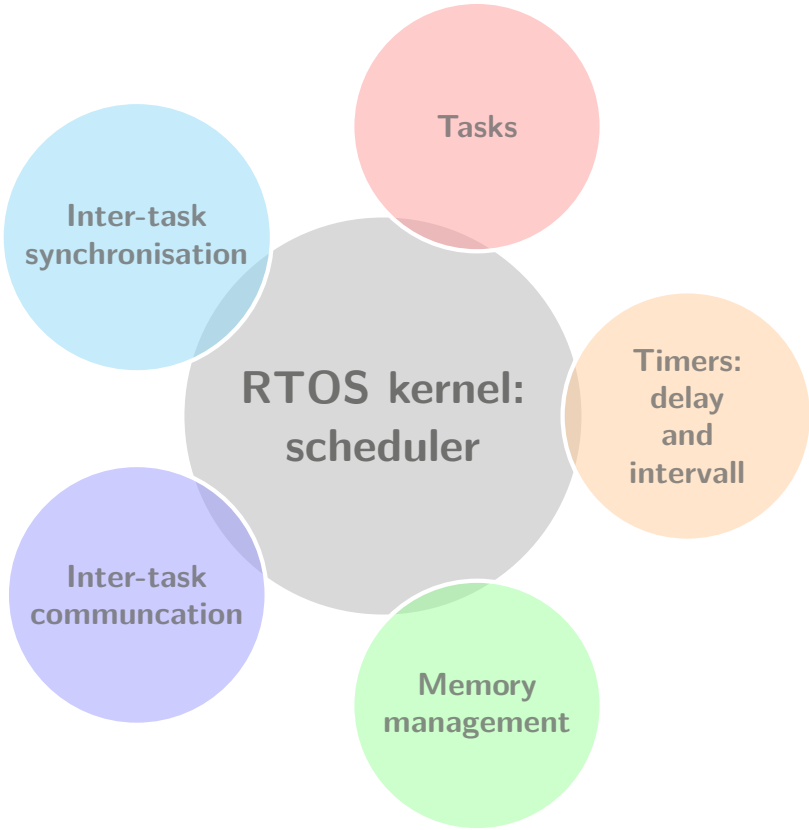
Intro

What is an RTOS?

Intro

Difference between
bare metal
and using
an **RTOS?**

Building blocks of an RTOS.



Advantages and disadvantages using an RTOS

Advantages:

- Allows more complex application designs
- Allows scheduling of tasks: priority based, time slicing, ...
- Perform multiple simultaneously running functions
- Better structured software
- Easier maintenance of software

Disadvantages:

- Larger footprint
- Requires additional knowledge
- May introduce some additional jitter or delay (SysTick + scheduler required)



Overview of different RTOS

- FreeRTOS
- Azure RTOS ThreadX
- VxWorks
- INTEGRITY RTOS
- Linux: Preempt RT (patch), RTLinux, Xenomai, L4Linux, ...
- ...

Some more: *Comparison of real-time operating systems* or *List of open source real-time operating systems*

Tasks

Tasks are the basic execution entities in real-time operating systems.

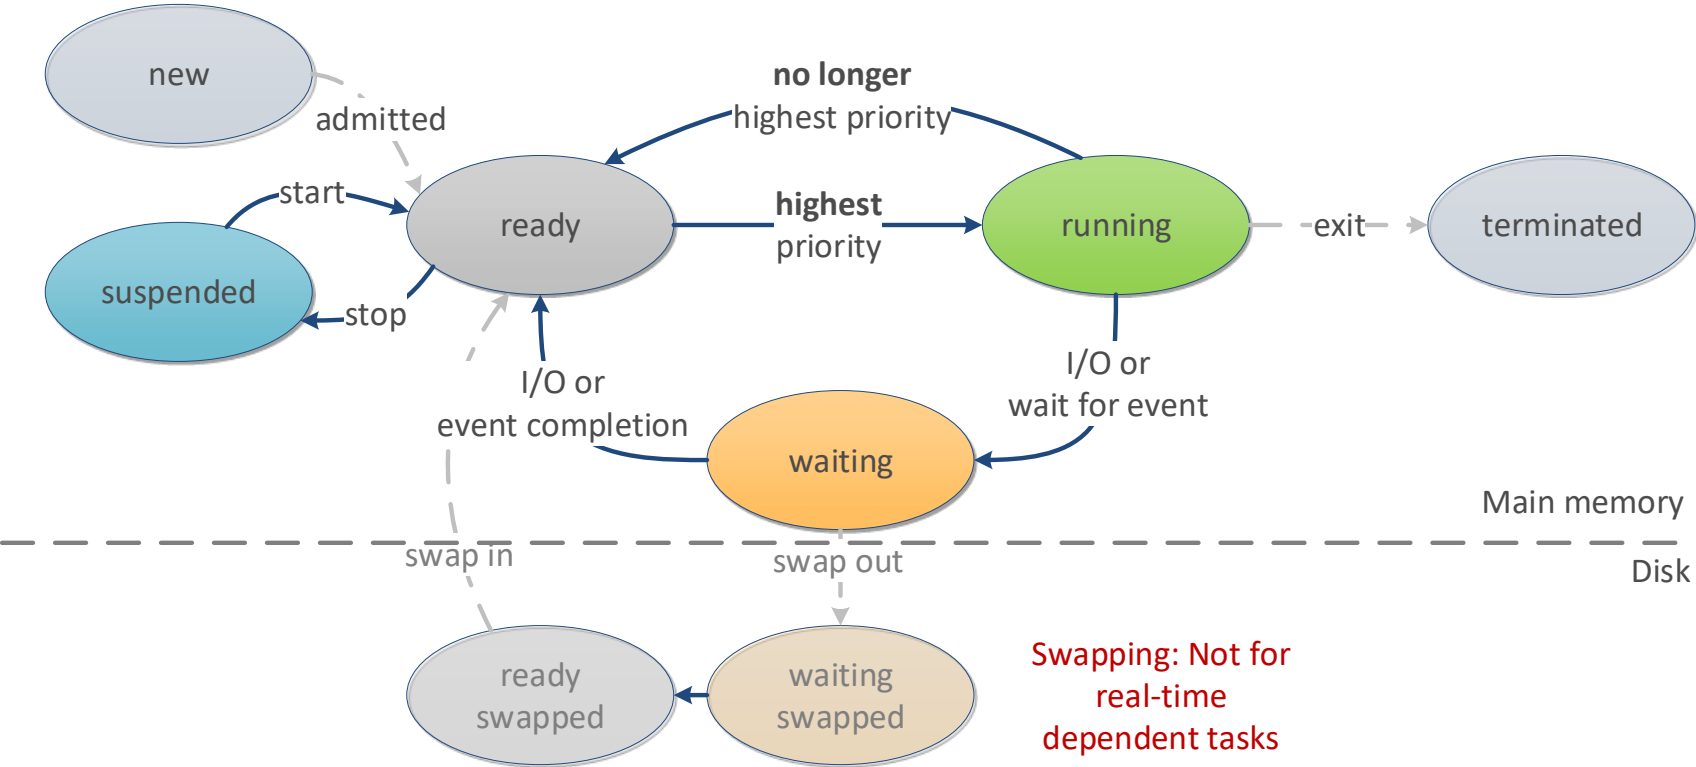
Operating system:

	Heavyweight process	Lightweight process
■ POSIX (original)	Process	Thread
■ Windows	Process: fork()	Thread: pthread_create()
■ Linux	Real-time process (RTP)	Task
■ VxWorks		Task
■ FreeRTOS		Thread
■ Azure RTOS ThreadX		

Real-time programming language:

■ Ada	Program, partition	Task
■ PEARL		Task

Task states



Synchronisation

To synchronise tasks and tasks with ISRs, different synchronisation mechanisms exist.

Mechanisms:

- Semaphore
 - Binary semaphore: often called Mutex
 - Counting semaphore
- Event flags

⇒ Some RTOS may offer additional synchronisation mechanisms (e.g. FreeRTOS: direct to task notification).

Semaphore

A semaphore is a signaling mechanism to control the access to common resources.

Operations:

Operation	Description
<code>seminit(s, value)</code>	Creates and initialises a semaphore with a value. The value is a number that specifies the number of processes that can simultaneously enter the critical area.
<code>P(s)</code>	Wait until the critical area is free (value--).
<code>V(s)</code>	Releases the critical area (value++).

Basic usage:

```

1 seminit(s, 1);
2
3 P(s);
4 //critical area..,
5 V(s);

```

Implementation:

Single-processor system:

- Disable interrupts (not preferred)

Multi-processor system:

- Test-and-set
- Compare-and-swap
- Spinlock (busy wait, for short waiting periods)

⇒ Can support: Polling, FIFO queueing, priority queueing, and priority inheritance.

Semaphore: Types

Types	Initialisation	Description
Mutex	<code>sem_init(s, 1)</code>	A mutex semaphore is used for mutual exclusion . Typically initialised with 1.
Binary	<code>sem_init(s, 0)</code>	A binary semaphore is used when there is only one shared resource . Initialisation with 0/1 possible.
Counting	<code>sem_init(s, N)</code>	A counting semaphore is used to handle more than one shared resource. Typically initialised with the number N of shared resources. Initialisation with 0/N possible.

Event flags

An event bit (or flag) is a one bit information to synchronise tasks or tasks with an ISR.

Often, the event bits are combined into an event group.

Usually these features are available:

- `EVENT_SET` Sets an event bit
- `EVENT_CLEAR` Clears an event bit
- `EVENT_WAIT` Waits until an event occurs; sometimes also OR, AND, NOT combinations possible

Additional possible:

- `EVENT_WAIT_ANY` Waits for any of the specified event bits (OR)
- `EVENT_WAIT_ALL` Waits for all of the specified event bits (AND)

Comparison of semaphores and events

Event flags

- State (bit information)
- On EVENT_WAIT, no change
- EVENT_WAIT of n tasks is satisfied by one EVENT_SET
- EVENT_SET and EVENT_CLEAR without an effect, if no task is waiting
- Typically: AND, OR available in EVENT_WAIT

Counting semaphores

Counter

- P() consumes an entity from counter
- P() of n tasks requires n V() operations
- V() (without an corresponding P()) increases the counter
- Typically: AND, OR **not** available in P()

Sometimes:

- Combination of semaphore and event flags possible
- Linux supports semaphore arrays which allow AND, OR

Communication overview

Different RTOS provide different communication mechanisms*:

Communication mechanism	Inter process	Over network	Stream (byte) oriented	Message oriented
■ Message queue	X			X
■ Stream buffers	X		X	
■ Sockets (TCP/IP)	x	X	X	
■ Pipes	X		X	x
■ Shared memory	X		X	x

*The concrete naming and the implemented feature set depends strongly on the RTOS and can differ significantly.

Communication

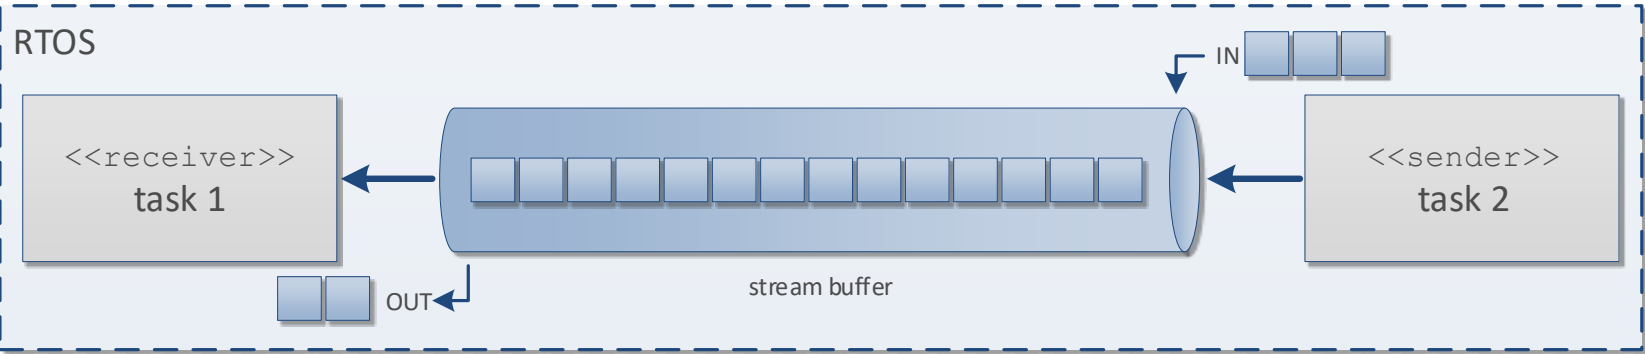
Message queues: a place for one or multiple messages



- Task to task communication
- Task to ISR communication

Communication

Stream buffers: a stream of bytes



- Task to task communication
- Task to ISR communication

Communication patterns

Common communication patterns:

- 1:1 communication (one queue, or multiple queues)
- n:1 communication (one queue, or multiple queues)
- 1:n communication (one queue, or multiple queues)
- m:n communication

Other variants:

- Messages buffered or unbuffered
- Blocking or non-blocking send/read
- Broadcast (send to n mailboxes in parallel)
- Communication via sockets (network, TCP/IP)


1:1 communication

1:1 communication with one message queue:



A fixed communication between two tasks.


1:1 communication with multiple message queue in a sequence:



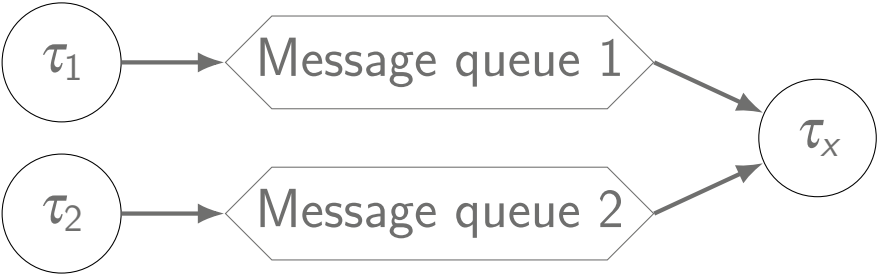
A more flexible communication pattern between two tasks.

n:1 communication

n:1 communication with one message queue




n:1 communication with multiple queues

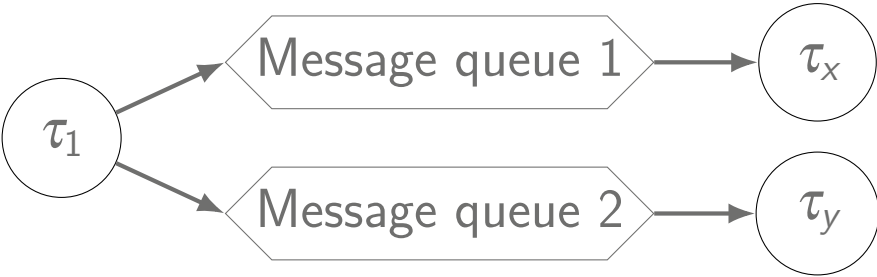


1:n communication

1:n communication with one message queue




1:n communication with multiple message queues



m:n communication

m:n communication with one message queue



Interrupts

Interrupts do have:

- No state information (like a TCB (task control block)):
 - No stack
 - No virtual PC
 - No own register set
- ⇒ Therefore, no blocking system calls are allowed (like `P()`)

But:

- Non-blocking system calls are possible (like `V()`)
- Secondary reaction: resume a task

Implementation:

- With assembler, or **with a higher level language: e.g. C/C++**
- Registration: directly within the interrupt vector table, or through an RTOS kernel wrapper function



Interrupts

Variant 1: Mini ISR

Mini ISR:

```
1 void ISR() {  
2     //e.g. transfer 1 word from/to I/O register  
3     //     to/from an internal buffer  
4     buffer[i] = ...  
5     ++i;  
6 }
```

Details:

- No interaction with the RTOS kernel (no system calls)
- Interruptions through such ISRs can usually always be allowed
- In practise usually not sufficient (because no notification is done)

Interrupts

Variant 2: ISR with system calls (non-blocking)

ISR with system calls (non-blocking):

```
1 void ISR() {  
2     ENTER_KERNEL(); //save registers, switch into system address space  
3  
4     //prepare message  
5     Message m = ...  
6     //e.g. MESSAGE_SEND m to task t  
7     MESSAGE_SEND(m, t);  
8 }
```

Details:

- Direct ENTER_KERNEL(); mechanisms required
- Such an ISR can not run during other kernel activities (e.g. MESSAGE_SEND();)
- Sometimes, ISRs are executed in the system address space (e.g. ARM Cortex M)

Interrupts

Variant 3: combination of 1 and 2

Mini ISR with system calls (non-blocking):

```
1 void ISR() {  
2     //e.g. transfer 1 word from/to I/O register to/from an internal buffer  
3     buffer[i] = ...  
4     if (buffer[i] == EOF){  
5         ENTER_KERNEL(); //save registers, switch into system address space  
6         V(TRANSFER_READY);  
7     }  
8     ++i;  
9 }
```

Details:

- Often used pattern
- Problem: When should such an ISR be allowed?

Interrupts

Variant 4: ISR with system call buffer

Mini ISR with system calls (non-blocking):

```

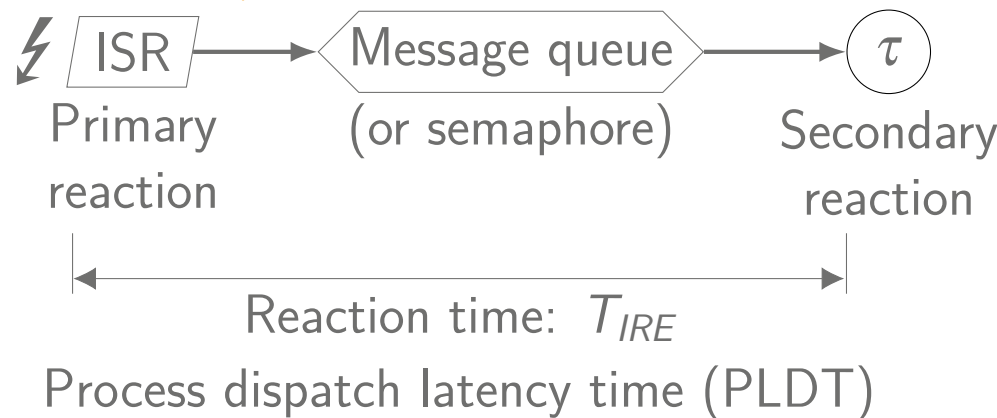
1 void ISR() {
2     //e.g. transfer 1 word from/to I/O register to/from an internal buffer
3     buffer[i] = ...
4     if (buffer[i] == EOF){
5         if(ENTER_KERNEL()) { //save registers, switch into system address space
6             V(TRANSFER_READY);
7         } else {
8             //Buffer: V(TRANSFER_READY);
9             //and execute it later through the RTOS
10        }
11    }
12    ++i;
13 }
```

Details:

- Also known as the deferred interrupt handling (or deferred procedure call in Windows)
- Minimises the execution time of ISRs (can improve response time for otherwise long running ISRs)

Interrupts

Primary/secondary reaction pattern



```
1 void ISR() {
2   Message m = ...
3   MESSAGE_SEND(QUEUE_ID, m);
4 }
```

```
1 void task() {
2   Message m;
3   //blocks until a message is available
4   MESSAGE_RECEIVE(QUEUE_ID, &m);
5
6   //work with message...
7   //can take a while
8 }
```

Timers

Typical time services in RTOS:

- Start a timer (timer task or function callback)
 - in x relative time units (e.g. in 5 seconds)
 - at absolute time x (e.g. at 12:23:01 o'clock)
 - periodic with cycle time x (e.g. 5 seconds with auto reload)
- Resume a timer (time specification as for start)
- Stop a timer
- Watchdog: Reporting a timeout (e.g. when waiting for I/O, semaphore, event, or message)

Implementation: HW vs SW timer

Hardware timer:

- Provided by the HW (timer or/and real-time clock)
- Usually, very accurate timing (low jitter)
- Requires the implementation of an ISR and a task notification mechanism

Software timer:

- Provided by the RTOS
- May be not as precise as HW timers (some jitter)



Trends

Multiprocessor or multi core systems are used for real-time systems

- n tasks run on m equally processors (CPU cores)
- Real parallel execution of tasks
- If correctly synchronised, everything works such as on a single core CPU
- Scheduling guarantees (and response times) are much harder to proof

Some details:

- Homogeneous vs heterogeneous multicore systems
- Symmetric vs asymmetric multiprocessing

Homogeneous vs heterogeneous multicore systems

Homogeneous multicore systems

- Multiple CPU cores
- Each CPU core has the same architecture
- Example: ARM Cortex A53 (quad core)

Heterogeneous multicore systems

- Multiple CPU cores
- Two or more CPU cores that differ in architecture
- Example: BeagleBone black (1 Cortex A8 (AM3358), 2x Cortex M3) or NXP i.MX7D
- Example: NXP i.MX7D (Cortex A7 + M4)

Advantages of heterogeneous systems:

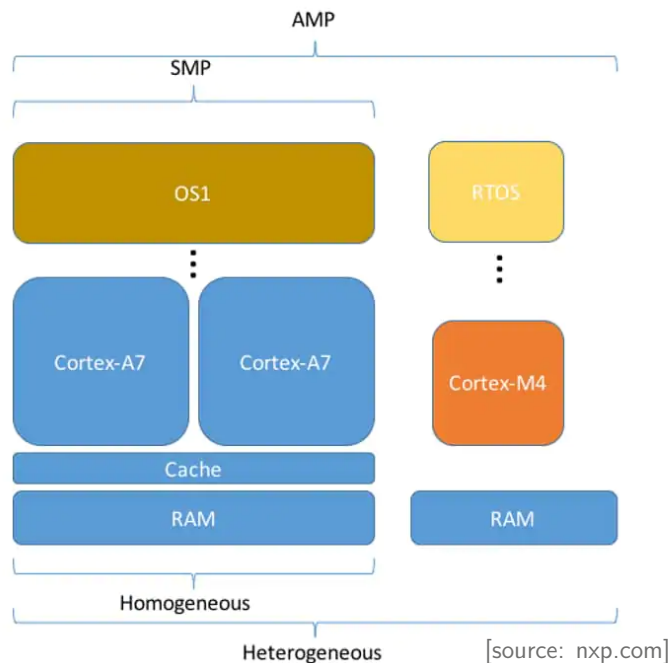
- Performance optimisation
- Reduction of power consumption
- Improved system reliability and security
- Processing real-time tasks and non real-time tasks (such as GUI) on the same system

[source: nxp.com]

Symmetric vs asymmetric multiprocessing

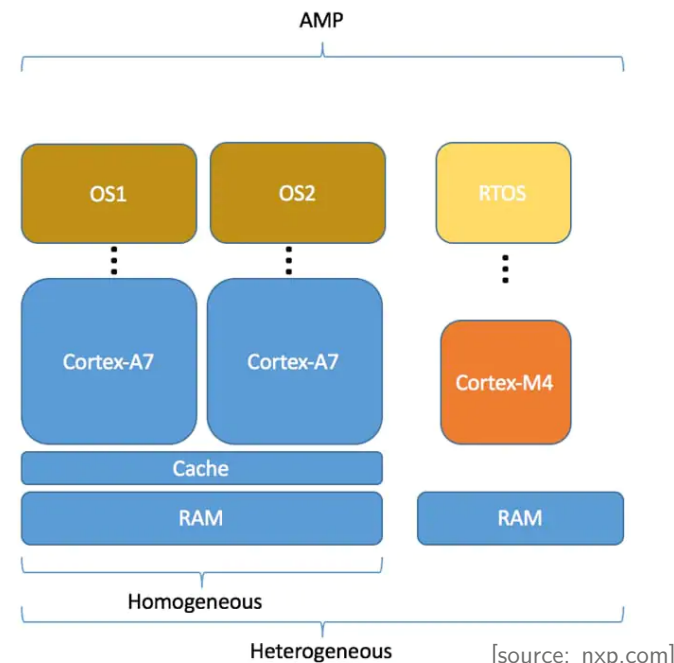
Symmetric multiprocessing (SMP)

- One kernel, multiple cores (usually identical cores)
- Enables load balancing and real parallel processing of tasks



Asymmetric multiprocessing (AMP)

- Contains multiple cores (homogeneous or heterogeneous cores)
- Usually, more than one OS is running



Summary and outlook

Summary

- Intro to RTOS
- Tasks
- Synchronisation
- Communication
- Interrupts
- Timers
- Trends

Outlook

- FreeRTOS